

Code Anomalies Flock Together

Exploring Code Anomaly Agglomerations for Locating Design Problems

William Oizumi, Alessandro Garcia,
Leonardo da Silva Sousa, Bruno Cafeo
Informatics Department – PUC-Rio
Rio de Janeiro, Brazil
{woizumi, afgarcia, lsousa,
bcafeo}@inf.puc-rio.br

Yixue Zhao
Dept. of Computer Science - USC
Los Angeles, USA
yixuezh@usc.edu

ABSTRACT

Design problems affect every software system. Diverse software systems have been discontinued or reengineered due to design problems. As design documentation is often informal or nonexistent, design problems need to be located in the source code. The main difficulty to identify a design problem in the implementation stems from the fact that such problem is often scattered through several program elements. Previous work assumed that code anomalies – popularly known as code smells – may provide sufficient hints about the location of a design problem. However, each code anomaly alone may represent only a partial embodiment of a design problem. In this paper, we hypothesize that code anomalies tend to “flock together” to realize a design problem. We analyze to what extent groups of inter-related code anomalies, named agglomerations, suffice to locate design problems. We analyze more than 2200 agglomerations found in seven software systems of different sizes and from different domains. Our analysis indicates that certain forms of agglomerations are consistent indicators of both congenital and evolutionary design problems, with accuracy often higher than 80%.

1. INTRODUCTION

Design problems are structures that indicate violations of key design principles or rules [38]. Every software system suffers from design problems, introduced either during original development or during evolution. Examples of design problems are *Fat Interfaces* [10, 38], *Overused Interfaces* [10, 38], and *Scattered Concerns* [10, 38]. These problems may have different degrees of severity, but all of them should be detected and possibly removed from the source code. Software systems have been often discontinued [22] or have had to be fundamentally reengineered [12, 34, 40] when design problems were allowed to persist in a system and to be compounded by other design problems introduced later.

Design problems are introduced and allowed to remain in a system because their localization in the source code is dif-

ficult. As design documentation is often informal or nonexistent, code anomalies – popularly known as code smells [7] – are used as surface indicators of design problems. Examples of code anomalies are *Long Method*, *Feature Envy*, and *God Class*. Even though each code anomaly can provide some hint to developers, it alone might not suffice to indicate the presence of a design problem. Each design problem is rarely localized in a single anomalous element; instead, it is scattered into different code elements of the implementation [30].

As an example, let us assume a developer is in charge of identifying *Fat Interfaces*. He will need to go through all code anomalies affecting each interface in the program. In particular, he will need to inspect all the classes that implement the interfaces in the source code. Then, he will have to analyze all the clients of such classes. Furthermore, it is hard and time-consuming to identify which code anomalies he should focus. Even for small software systems, there are hundreds of code anomalies [25] and thousands of possible relationships to examine.

The relation between design problems and their counterpart code anomalies is often complex. Unfortunately, there is no understanding of which relationships between code anomalies are frequent indicators of design problems in an evolving program. There is a recent growing interest in conceptually characterizing interactions between code anomalies [1, 2, 30, 36, 44]. However, the relation of code anomalies and design problems is rarely investigated. Empirical studies only address how individual occurrences of code anomalies emerge during software evolution [39] or affect quality attributes [5, 13, 17, 18, 21, 32]. They do not analyze how individual anomalies and their relationships in the code might help developers to spot design problems.

In this paper, we hypothesize that code anomalies “flock together” in order to embody different design problems in the implementation. Therefore, we investigate whether and how code anomaly relationships can help developers to locate design problems. To achieve this purpose, we propose a strategy to identify groups of inter-related code anomalies, i.e., code anomalies that “flock together”. We call *agglomerations* these groups of inter-related code anomalies. We perform a multi-case study in order to investigate them, focusing on *agglomerations* of code anomalies that are syntactically or semantically related (Sections 3.2 and 3.3). We analyze seven systems of different sizes (8 KSLOC to 129 KSLOC) and from different domains. Our analysis involves a total of 5418 code anomalies and 2224 agglomerations.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ICSE '16, May 14–22, 2016, Austin, TX, USA

© 2016 ACM. ISBN 978-1-4503-3900-1/16/05...\$15.00

DOI: <http://dx.doi.org/10.1145/2884781.2884868>

We investigate the circumstances under which agglomerations are related (or not) to design problems. The analyzed circumstances involve: (i) the statistical significance of the relation between agglomeration types and design problems, (ii) the strength of this relation, and (iii) the extent to which agglomerations manifest themselves at different stages of a system’s lifetime, i.e., in early vs. late versions of a system. The aforementioned analysis results in several findings:

1. Overall, approximately 50% of the syntactic agglomerations are related to design problems. Their use can help developers to discard almost 4000 non-agglomerated code anomalies that are irrelevant to locate design problems. In general, at least 3 code anomalies of each syntactic agglomeration are related to the same design problem in all the systems.
2. However, syntactic agglomerations do not suffice to assist developers in locating all design problems. Semantic agglomerations are much more consistent indicators of design problems across our target systems. On average, 80% of the semantic agglomerations are related to design problems. Several design problems in all the analyzed systems can only be revealed with semantic agglomerations.
3. Analysis of history co-changes has been recently used to identify design problems during software evolution [35, 41, 42]. However, we observe this approach is ineffective in locating a significant proportion of design problems. The reason is that those problems are “congenital”, i.e., they were already introduced in the system’s initial version. Co-changes affecting the corresponding anomalous elements might occur only in later versions of a software system, but it is usually too late to identify and remove congenital design problems. Many clients already depend on the anomalous code elements that realize the design problem.

The paper is organized as follow. Section 2 presents the background to understand the paper. Section 3 describes the settings of our study, including the research questions and the procedure for data collection and analysis. Section 4 summarizes the main results of our study. Sections 5 and 6 present the related work and the threats to validity, respectively. Finally, Section 7 concludes the paper.

2. BACKGROUND

This section presents the background required to understand our study settings. Section 2.1 presents key concepts, and Section 2.2 introduces a motivating example.

2.1 Code Anomalies and Design Problems

A code anomaly, popularly known as “code smell”, is a micro structure in the program that represents a surface indication of a design problem [7]. Examples of code anomaly types vary from method-level smells, such as *Long Method* and *Feature Envy*, to class-level smells, such as *God Class*, *Data Class*, *Shotgun Surgery*, and *Divergent Change* [7, 19]. Even in small programs, developers often have to face hundreds or thousands of code anomaly instances. They need to analyze all these anomalies in order to discard, postpone or further consider them. In fact, multiple anomalies may

interact in the source code structure, at which each one possibly represents a partial hint of a design problem. On the other hand, many anomalies may not contribute at all in helping developers to locate any design problem [23, 24, 25].

An important set of code anomalies actually represents design problems. Design problems are structures that indicate violation of intended design rules or fundamental design principles, and they negatively impact design quality [8, 38]. Examples of relevant design problems are *Fat Interface* [26] and *Unwanted Dependencies* [33]. The former is a general, ambiguous entry point of a design component that provides non-cohesive services, thereby complicating the logic of its clients. This design problem violates the well-known principles of cohesion, abstraction, and separation of concerns. The latter represents the violation of a design rule that specifies that two design components should not communicate.

In this study, we focus on such design problems that affect a system’s design decomposition into major sub-systems (components) and their interfaces. Such design problems are often targets of major maintenance efforts [10, 34, 43], and they are very often associated with design degradation that leads to partial or full discontinuation of a software system [12, 22, 34, 40]. Therefore, they should be removed as early as possible from a system.

Design problems often affect multiple code elements and are hard to identify in the source code. Then, code anomalies analyzed individually may not be sufficient to help developers in diagnosing a design problem. Additionally, certain anomaly types convey micro-structures that are the natural implementation solutions given the role of the code element. For instance, some types of classes (e.g., lexical-analyzer classes) naturally address several concerns, while some methods (e.g., command processing and dispatching) are naturally long [19]. Thus, certain code anomalies neither cause nor contribute to the location of a design problem.

2.2 Motivating Example

Let us consider the example shown in Figure 1 extracted from Apache OODT (Object Oriented Data Technology) [27]. OODT is a data-grid system aimed at supporting the management and storage of scientific data. OODT’s *Versioning* component, which exports the *Versioner* interface, is responsible for managing and storing versions of different *Product* types using different storage strategies. All classes implementing *Versioner* interface have to implement the *createDataStoreReferences* method. One of the parameters of this method is a *Product* instance. *Product* class can represent several types of structures, like flat and hierarchical. As there are no classes for each *Product* type, each *createDataStoreReferences* implementation has to decide if it is dealing with the right *Product* type. For example, the *MetadataBasedFileVersioner* class only deals with “flat” products.

Over time, OODT developers realized that there may be a design problem located in the *Versioning* component. Scattered changes involving this sub-system have been a major source of maintenance effort. They ran a static analysis in order to identify code anomalies that point to the design problem. The analysis output consisted of several dozens of code anomalies scattered across several classes and methods of the *Versioning* component. They analyzed several anomalies individually and decided to discard some of them, as they did not represent a threat to the program structure.

Finally, developers zero in on a frequently changed method of *BasicVersioner*, called *createDataStoreReferences*. In fact, this method is simultaneously affected by four anomalies: *Long Method*, *Feature Envy*, *Shotgun Surgery* and *Divergent Change*. However, such local analysis does not suffice to conclude whether this is part of a major design problem. OODT developers begin looking for other anomalies at neighboring methods and classes, i.e., those methods and classes that have syntactic relationships with *createDataStoreReferences* and *BasicVersioner*. They work their way “outward” by navigating in the hierarchy structure, and they move up to analyzing all clients of the *Versioner* interface, such as *GenericFileMngObjFactory* and *XmlRpcFileMngrClient*.

After observing all the anomalies affecting the hierarchical structure of *Versioning* and its direct clients, the developers are able to infer that the component is being affected by the *Fat Interface* problem: the *Versioner* interface is a single entry point of the component, but only the anomalies (e.g., *Feature Envy*) scattered in its four implementations reveal the non-cohesive nature of the interface. Those classes need to change every time that each of the different product types are changed (i.e., *Divergence Changes*). The scattered occurrences of *Shotgun Surgeries* in the *Versioning* component’s clients reinforce the interface is providing several non-cohesive services, which should be segregated into separate interfaces. Each client also needs to change every time the list of products and their details are changed. All the components that are related to *Versioner* and have both anomalies (*Feature Envy* and *Shotgun Surgery*) are inside of the dashed line in the Figure 1. These are the components affected by *Fat Interface*.

The presence of *Fat Interface* in the *Versioning* component is causing other problems. While a *Fat Interface* decouples components, it makes the component less understandable and analyzable. Determining the actual services exposed by such component requires inspecting its implementation details. Furthermore, the generality of the interface also makes it easier to misuse, since different functionalities are exposed by the same interface. Yet the only way this potentially critical design flaw can be discovered is by reflecting upon multiple related code anomalies that are located in syntactically related modules: the interface, its subclasses, and the interface clients.

3. STUDY DEFINITION

Our study investigates whether and how inter-related code anomalies, referred to *anomaly agglomerations*, are related to design problems. Section 3.1 describes and motivates our research questions. Sections 3.2 and 3.3 present the types of agglomerations we are investigating. Section 3.4 describes the target systems of our study. Finally, Section 3.5 addresses the procedure for data collection and analysis.

3.1 Research Questions

Existing techniques [1, 30, 36] usually assume that individual anomalies suffice for assisting developers in locating design problems in the program. The previous section showed that each design problem may be realized by various inter-related code anomalies scattered in the program. However, there is no understanding whether certain anomaly relationships can help developers in better locating design problems than individual code anomalies. We are primarily interested in groups of syntactically related code anomalies,

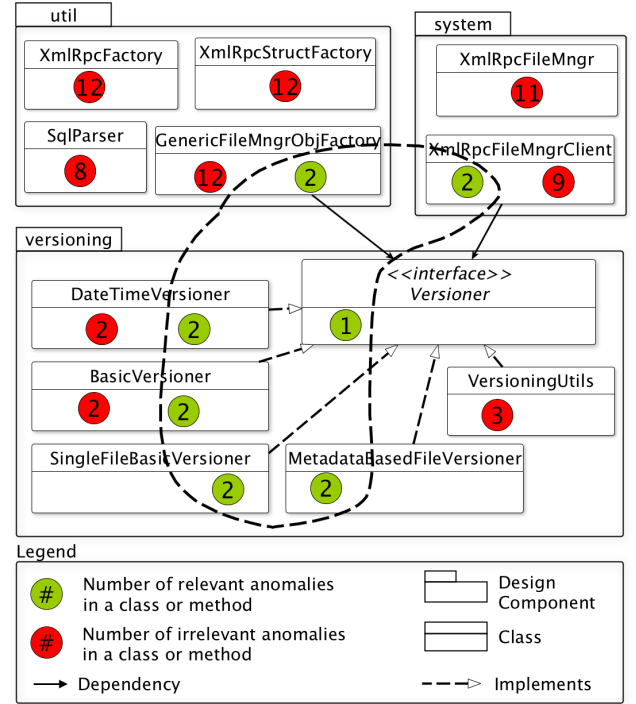


Figure 1: Fat interface affecting Versioner implementation

i.e., *syntactic agglomerations* for short (Section 3.2). Two anomalies are syntactically related if their host program elements are connected through method calls or inheritance relationships. We first investigate whether syntactic agglomerations often embody design problems in the source code. If this hypothesis holds, the inspection of such agglomerations can help developers in locating design problems. This reasoning leads to our first research question:

RQ1 Are syntactic anomaly agglomerations sufficient indicators of design problems?

Syntactic agglomerations may not suffice to locate all design problems. Thus, other forms of code anomaly agglomerations may be required to locate design problems in the source code (Section 3.3). Therefore, we address RQ1 by analyzing the relation between diverse forms of agglomerations and design problems in the context of seven software projects. We also analyze this relation in early versions of each project. Recent studies [2, 39] revealed that, most of the time, programs are affected by code anomalies since their creation. In addition, it might be that some design problems are “congenital”, i.e., they manifest in the initial versions of a program (Section 2.2). However, there is no understanding about the relation between “early” code anomalies and congenital design problems; and more importantly, between anomaly agglomerations and design problems. This gap motivates our second research question:

RQ2 What proportion of design problems manifest as anomaly agglomerations in early versions of a program?

We address RQ2 by investigating the relation of design problems and code anomaly agglomerations in early versions of a system. If this relation holds, early detection of congenital design problems can be improved by using code anomaly

agglomerations. We address RQ2 by analyzing the available initial versions of the analyzed software systems.

3.2 Syntactical Agglomerations

A *syntactic agglomeration* is a group of at least two anomalous code elements explicitly related in a program. A code element is *anomalous* when it is affected by one or more code anomalies. In our study, we consider the following code elements: classes, interfaces, methods, constructors, and fields. An explicit relationship is established between two anomalous elements when they have at least one of the following dependencies: shared attribute, method call, class extension or method overload. We classify syntactic agglomerations according to their scope in the program: (i) *intra-component agglomerations*, i.e., those entirely located within a single component, and (ii) *inter-component agglomerations*, i.e., those located in two or more components. In our study, we consider each component is realized by a package in the analyzed Java programs (Section 3.4), unless the system developers specify otherwise. Developers may state a component is realized by a set of classes that are not necessarily located in the same package.

Inter-Component Agglomeration is a syntactic agglomeration that involves two or more design components, with at least one anomalous code element located within each of them. The example in Section 2.2 illustrates an inter-component agglomeration in the OODT system. The agglomeration is formed by anomalous code elements located in three OODT components: *Versioning*, *Util* and *System*. In Figure 1, the inter-component agglomeration is surrounded by the dashed black line, and the green circles represent the code anomalies located in the syntactically related code elements that compose the agglomeration.

Listing 1 illustrates the algorithm (pseudo-code) for computing inter-component agglomerations. The algorithm has two parameters: (i) *dc*: a set of design components, and (ii) *agglomThreshold*: a threshold value for the minimum agglomeration size. First, the algorithm identifies all anomalous elements of a design component by using the function *getAnomalousElemsOfComp()* (line 6). The anomalous elements of the program are recorded in the *anomalousElems* attribute. The algorithm generates a graph (*graphInterElems*) with the syntactic relationships (edges) between anomalous elements (vertices) located in different design components (line 10 and lines 12–19). To find the relationships between elements, the algorithm uses the function *isRelated()* (line 15). In order to detect the agglomerations, the algorithm finds the subgraphs within the graph in which any two vertices are connected (the subgraph formed by the relationships among the anomalous elements). Finally, only the agglomerations with the size greater than *agglomThreshold* are included in the resulting set of inter-component agglomerations, i.e., *interAgglomerations* (lines 25–29).

Intra-Component Agglomeration is a syntactic agglomeration composed of anomalous code elements located in the same design component. Given space constraints, the algorithm is available in our on-line supplementary material [31]. However, the algorithm for computing intra-component agglomerations is similar to the previous algorithm for identifying inter-component agglomerations. The only difference is that the algorithm generates a graph with the relationships (edges) between the anomalous elements (vertices) within each design component.

Listing 1: Inter-component Agglomeration Algorithm

```

1. public List<> getInterAgglomeration(Set dc, int
   agglomThreshold){
2. InterAgglomerations = {}
3. anomalousElems = {}
4.
5. for(each design component dc in the program){
6.   anomalousElems.addAll(getAnomalousElemsOfComp(dc)
   )
7. }
8.
9. //Adding anomalous code elements as vertices in
   the graph
10. graphInterElems = new graphInterElems().addVertex
   .addAll(anomalousElems)
11.
12. for(each e1 in anomalousElems){
13.   for(each e2 in anomalousElems){
14.     //Creating the Inter-component graph
15.     if(isRelated(e1,e2) && (e1.getDesignComponent
   () != e2.getDesignComponent())){
16.       graphInterElems.addEdge(e1,e2)
17.     }
18.   }
19. }
20.
21. // Adding the connected components as
   agglomerations
22. InterAgg = new InterAgg().addAll(
   getConnectedComponents(graphInterElems))
23.
24. // Selecting only agglomerations above the
   predefined threshold
25. for(each a in InterAgg){
26.   if(size(a.getAnomalies()) > agglomThreshold){
27.     InterAgglomerations.add(a)
28.   }
29. }
30. return InterAgglomerations

```

3.3 Semantic Agglomerations

Based on the example of Section 2.2, we hypothesized that code anomalies might somehow interact through the program structure because of the presence of a single design problem. In that example, the syntactic relationships amongst the anomalous elements would be sufficient to help developers in revealing the design problem. However, this might not be the case for all occurrences of design problems. In some cases, design problems might be related to semantically connected anomalies. Anomalies are semantically connected if their host elements are addressing the same concern. Concern is a property or functionality of interest to the designers of a system, but it is not necessarily modularized in a single component. In such cases, the semantic relationship may help the developers to locate certain design problems better than the syntactic relationships.

Let us consider an example extracted from the Mobile Media system [45]. Mobile Media is a software product line to derive applications that manipulate photos, videos and music on mobile devices [45]. Figure 2 depicts a partial view of three components of Mobile Media design: *Controller*, *Screens*, and *Sms*. Classes of the *Screens* component are affected by the *Divergent Change* anomaly, while classes of the *Controller* and *Sms* components were infected by the *Divergent Change* and *Shotgun Surgery* anomalies. These code anomalies did not represent isolate problems. Many of these classes are not syntactically connected, but their methods partially address the same concern, called *Photo Label Management*. Therefore, the anomalous code elements are altogether realizing the design problem *Scattered Concern*, i.e., multiple components realizing a crosscutting concern

[9, 24]. The realization of *Photo Label Management* should be modularized in the *Controller* component. This problem was the cause of major design refactorings along Mobile Media evolution. However, the design problem would be better spotted if the scattered anomalies (*Divergent Changes* and *Shotgun Surgeries*) realizing the same concern are considered altogether. In our study, we chose concern-based agglomeration as a representative type of semantic agglomerations.

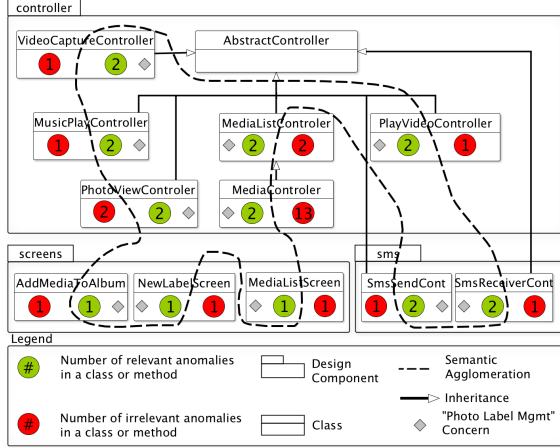


Figure 2: Scattered Concern in Mobile Media

Semantic Agglomeration is composed of anomalous code elements realizing a single concern, which is not modularized by design components. Examples of design-relevant concerns can be classified in domain-specific concerns, such as *Photo Label Management*, or general-purpose concerns, such as persistence, error handling, security and the like. Semantic relationships between two or more code anomalies occur when their host code elements are intended to (partially) realize a single design’s purpose or concern, i.e., the semantic relationship appears in code elements that implement the same concern and also have code anomalies.

Listing 2 presents the algorithm (pseudo-code) for detecting semantic agglomerations. The full algorithm is available in our supplementary material [31]. The basic idea of the algorithm is searching for agglomerations for each concern *con* that satisfies the following conditions: (i) *con* is located in one or more components (line 7), and (ii) at least one of these components is not mainly dedicated to realize *con* (line 5). The components in (ii) are named *weakly-dedicated components*. The identification of weakly-dedicated components is implemented by the function *weakDedicatedComponents()* (line 5). This function computes, for each component realizing a concern, if code elements in the component are mostly dedicated to realize other concerns. In other words, a component *dc* is weakly dedicated to a concern *con*, if *con* is not the main concern of the code elements within *dc*. Although *dc* partially implements *con*, the predominant concern of *dc* is not *con*. The minimum degree of dedication is captured based on a percentage threshold, named *weakThreshold*.

Thus, the main algorithm uses four inputs: a set of design components *dc*, the mappings for each concern *con*, a threshold for the minimum agglomeration’s size *agglomThreshold* (default = 1), and *weakThreshold* (default = 1). In our case study, the developers provided the concern mappings (Section 3.5), but the concerns can be automatically extracted from the source code using an automatic concern location

Table 1: Characteristics of the Target Systems

System	Type	Design	KSLOC
HW	Web Framework	Layers	8
MM	Software Product Line	MVC	10
S1	Desktop Application	Client-Server	122
S2	Desktop Application	Client-Server	118
S3	Desktop Application	Client-Server	93
S4	Web Application	MVC	116
OODT	Middleware	Layers	129

tool. Once the concerns are recovered, the algorithm can locate the code elements that implement an concern *con* and also are infected by code anomalies (line 7). The anomalous code elements contributing to weakly-dedicated components of a crosscutting concern are grouped into an agglomeration candidate (lines 8 to 9). If the agglomeration candidate has a number of anomalies higher than the *agglomThreshold* (line 11), then this agglomeration is included in the results (line 12) and confirmed as an actual agglomeration.

Listing 2: Semantic Agglomeration Algorithm

```

1. public List[] getSemanticAgglomeration(Set dc, Set
   con, int agglomThreshold, int weakThreshold){
2.   semanticAgglomerations = []
3.
4.   for(each design concern con in the program){
5.     W = weakDedicatedComponents(con, weakThreshold)
6.     agglomeration = []
7.     if(size(W) > 0){
8.       ae = getAnomalousElementsPerConcern(W, con)
9.       agglomeration.addAll(ae)
10.    }
11.    if(size(agglomeration) > agglomThreshold){
12.      semanticAgglomerations.add(agglomeration)
13.    }
14.  }
15.  return semanticAgglomerations
16. }

```

3.4 Target Systems

In order to address the two research questions, we analyzed systems with a wide range of characteristics. We selected 7 systems of different sizes, leveraging different design styles, and spanning different domains. Table 1 summarizes the general characteristics of each system. We focused on these systems because: (i) their designs had degraded, (ii) they present a wide range of design problems, and (iii) their developers were available to provide us with a reliable list of design problems (which are causes of major maintenance effort) and the mappings of their design’s concerns.

The first system is the Health Watcher (HW), a web framework system that allows citizens to register complaints about the health issues in public institutions [37]. The second is Mobile Media (MM), an academic software product line to derive applications that manipulate photos, videos, and music on mobile devices [45]. The next four systems are proprietary and, due to intellectual-property constraints, we will refer to them as S1, S2, S3 and S4. S1 and S2 manage activities related to production and distribution of oil. S3 manages the trading stock of oil, and S4 supports the financial market analysis. Finally, the last system is Apache OODT, whose goal is to develop and promote the management and storage of scientific data [27]. For all the target systems, several classes implement each component. In OODT, for instance, each design component is implemented by an average of 24 classes. There was always a 1-to-1 map-

Table 2: Analyzed Design Problems

Name	Description	Instances
Fat Interface	Interface of a design component that offers only a general, ambiguous entry-point that provides non-cohesive services, thereby complicating the clients' logic.	114
Unwanted Dependency	Dependency that violates an intended design rule.	2145
Component Overload	Design components that fulfill too many responsibilities.	141
Cyclic Dependency	Two or more design components that directly or indirectly depend on each other.	351
Delegating Abstraction	An abstraction that exists only for passing messages from one abstraction to another.	35
Scattered Concern	Multiple components that are responsible for realizing a crosscutting concern.	216
Overused Interface	Interface that is overloaded with many clients accessing it. That is, an interface with 'too many clients'.	39
Unused Abstraction	Design abstraction that is either unreachable or never used in the system.	59

ping between components and packages in three systems: MM, HW, OODT, S3 and S4. For S1 and S2, the developers provided the set of classes implementing each component.

3.5 Data Collection Procedure

The data collection process comprised the following activities: (i) identifying design problems with help of the systems developers, (ii) detecting code anomalies, concern mappings and agglomerations, and (iii) correlating agglomerations and design problems. Next, we describe each activity.

Identifying Design Problems. We produced a “ground truth” of design problems for each target system. We performed two steps to incrementally develop the ground truth. First, original developers of the systems provided us with an initial list of design problems. They listed the problems and explained the relevance of each one through a questionnaire [31]. They reported which was the maintenance effort caused by the presence of each design problem. They also described which code elements were contributing to the realization of each design problem. Second, we performed other steps to validate the initial list for correctness and completeness. An additional identification of design problems was performed using the source code and the system design. For systems without design documentation, we relied on a suite of design recovery tools [9]. The procedure for deriving the list of design problems with developers was the following: (i) an initial list of design problems was identified using detection strategies presented in [2], (ii) the developers had to confirm, refute or expand the list of identified design problems, (iii) the developers provided a brief explanation about the relevance of the design problem, and (iv) when we suspected there was still inaccuracies in the list of design problems, we asked the developers for further feedback. Table 2 describes the types of design problems and number of instances identified in our sample of systems.

Concern Mappings. The initial lists of concerns and their mappings in the source code were provided with the assistance of systems’ developers. For each concern in their initial lists, they provided a list of methods or classes realizing those concerns. Given the large size of certain systems, developers could eventually not produce complete concern mappings. Therefore, we also relied on the Mallet [28] tool. Mallet is a concern location tool that explores topic modeling in order to automatically generate a list of concerns and identify the code elements realizing each concern in a program. Then, we computed and compared two sets of semantic agglomerations: one based on the concern mapping produced by developers, and the other based on the concern mapping generated by Mallet. We will discuss the comparison of such agglomerations in Section 4.2.

Code Anomalies and Agglomerations. We considered 7 types of code anomalies: *Data Class*, *Divergent Change*, *God Class*, *Shotgun Surgery*, *Feature Envy*, *Long Method* and *Long Parameter List*. They were selected because they represent different types of symptoms at class and method levels. The detection of code anomalies was performed using well-known metrics-based strategies [19]. These detection strategies are similar to those typically used in previous empirical studies (e.g. [23, 24, 25, 39]). Such strategies have been proven to be effective for detecting code anomalies in other systems, with precision higher than 80% [2, 24]. Once the anomalies were detected, we computed the syntactic and semantic agglomerations. The detection of anomalies and agglomerations were carried out with a tool, called Organic [23]. Organic implements: (i) the 7 detection strategies for code anomalies, and (ii) the algorithms for computing agglomerations (Section 3.2 and Section 3.3). The results of these steps are available at the supplementary material [31].

Correlating Agglomerations and Design Problems.

As aforementioned, developers indicated the location of each design problem. In order to answer our research questions, we defined a criterion for correlating design problems with anomalies and agglomerations. A design problem is related to an individual code anomaly if the former is either partially or fully realized by the code element affected by the anomaly. We consider that a design problem and an agglomeration are related if they co-occur in at least two code elements. Even though agglomeration and design problem may be located in more than two anomalous elements, our criterion considers sufficient if they co-occur in two elements. Thus, an agglomeration fails to indicate design problems when either none or only one of its code anomalies is related to a design problem. Finally, we computed the strength of the relation between agglomerations and design problems, i.e., the number of anomalies in agglomeration that contributed (or not) to a design problem (Section 4.1). We checked the statistical significance of our results using the Fisher’s exact test [6] and computing the Odds Ratio [4] with the R tool [3].

4. RESULTS AND ANALYSIS

This section presents the results and analysis of our study. Section 4.1 discusses whether syntactic agglomerations assist the location of design problems. Then, Section 4.2 analyzes to what extent syntactic agglomerations suffice to locate all design problems and, if not, whether semantic agglomerations can further improve this task. Section 4.3 discusses if agglomerations can also be effective to assist the location of design problems in early versions of a system.

4.1 Exploring Syntactic Agglomerations

Before answering our first research question, we reflect upon the relation of syntactic agglomerations and design problems. First, we check whether syntactic agglomerations are more related to design problems than non-agglomerated code anomalies. Non-agglomerated anomalies are “detached” anomalies that do not take part of any agglomeration.

The results are presented in the first three columns of Table 3. The first column (named Agglomeration) lists the agglomeration types analyzed in our study. The second column (named AG-DP) shows the amount of agglomerations (for each category) related to design problems. The third column (named NoAG-NoDP) presents the number of “detached” anomalies that are irrelevant to locate design prob-

Table 3: Contingency Table and Fisher’s Test Results

Agglomeration	AG-DP	NoAG-NoDP	p-value	ORs
Intra-component	247	3996	<0.001	5,669636
Inter-component	167	4207	<0.001	4,878982
Syntactic	296	3759	<0.001	5,513531
Semantic	97	4463	<0.001	7,392637
All	312	4596	<0.001	2,999686

lems, i.e. they are not related to any design problem. We discuss the data of the last two columns later in this section. The data of the row “semantic” is discussed in Section 4.2.

Reducing the Search Space for Design Problems. Table 3 confirms that syntactic agglomerations are good indicators of design problems. Almost 300 syntactic agglomerations are related to design problems. Each row shows how many “detached” anomalies could be discarded when developers are exploring a specific type of syntactic agglomeration in order to look for design problems. For instance, if developers are inspecting inter-component agglomerations, they will be able to: (i) find 167 design problems in their systems, and (ii) discard more than 4000 code anomalies returned by their static analysis tools. This result suggests that syntactic agglomerations can assist developers in locating design problems. When identifying design problems in the source code, the scope of analysis can be reduced to the list of code anomalies taking part in the agglomeration, rather than an unmanageable list of individual code anomalies.

Syntactic Agglomerations as Design Problems: Strength of the Relation. Even though the previous analysis is interesting, it does not consider that anomaly agglomerations have higher probability of being related to design problems than a single code anomaly, since agglomeration involves more code. The latter affects a single code element (a method or a class), while the former involves at least two code elements. The size of an agglomeration ranged from 2 to 9 in all systems. Therefore, we investigated the strength of the relation between agglomerations and design problems. We performed further analyses to check to what extent various code elements of the agglomeration are, in fact, contributing to the realization of a design problem.

First, we calculated the Odds Ratio [4] for each type of agglomeration (ORs column of Table 3), which shows how much higher is the chance of code elements taking part in agglomerations to be related to design problems than individual anomalous elements. The chance of each anomaly within a syntactic agglomeration being related to a design problem is more than five times (5,513) higher than each “detached” anomaly. Similar results are observed for both intra-component and inter-component syntactic agglomerations. Second, we also observe that almost all syntactic agglomerations had at least 3 anomalous code elements simultaneously related to the same single design problem in most of the systems. Thus, when fixing a particular design problem, developers are more likely to reveal several anomalous code elements involved in the refactoring strategy.

Syntactic Agglomerations as Design Problems: Statistical Significance. Finally, we used the Fisher’s exact test to analyze the statistical significance of the relation between syntactic agglomerations and design problems. The fourth column (p-value) in Table 3 shows this correlation for each agglomeration type. We assume a confidence level higher than 99% (p-value threshold = 0.001) as the threshold value to the significant level of the test. Applying the

test, we observe that for all agglomeration types the p-value was lower than 0.001. There is a high correlation between both forms of syntactic agglomerations and design problems in the target systems. However, it does not necessarily mean that syntactic agglomerations should not be complemented with other forms of agglomerations to further improve the location of design problems. Therefore, we explicitly address our first research question in the next section.

4.2 Do Syntactic Agglomerations Suffice?

Even though syntactic agglomerations are statistically related to design problems, it is important to understand if syntactic agglomerations alone suffice to locate all design problems. Thus, this section addresses RQ1: “Are syntactic anomaly agglomerations sufficient indicators of design problems?” In order to answer RQ1, we first analyzed the proportion of syntactic and semantic agglomerations (un)related to design problems for each target system. We compared how often syntactic and semantic agglomerations relate to design problems. Table 4 shows, for each agglomeration type, the following measures: (i) **Precision**, (ii) **Recall**, and (iii) absolute number of design problems related to agglomerations (**DP**). Moreover, for each agglomeration type, Table 4 shows the median and the standard deviation for the columns **P** and **R**. The numbers between parentheses in the **DP** columns are discussed later in this section.

Precision indicates the fraction of identified agglomerations that are correctly related to design problems. Recall indicates the fraction of design problems successfully identified by agglomerations. We calculate precision and recall measures using the following standard equations:

$$P = \frac{TP}{TP + FP} \quad (1) \quad R = \frac{TP}{TP + FN} \quad (2)$$

where, TP (true positive) and FP (false positive) encompass all agglomerations that, respectively, are or are not related to design problems. FN (false negative) occurs when a group of code elements is affected by a design problem, but none of them is part of an agglomeration.

Table 4 shows the relation between agglomerations and design problems in terms of precision and recall. The Syntactic column indicates the aggregate results for both types of syntactic agglomerations: intra-component (second column) and inter-component (third column). We highlight that the values of the Syntactic column cannot be directly obtained by the data in the previous columns. In other words, the number of design problems in the fourth column is not the result of summing the numbers of design problems for both types of syntactic agglomerations. The reason is that there are intersections between instances of intra- and inter-component agglomerations. That is, an inter-component agglomeration might be composed of anomalies that are also members of one or more intra-component agglomerations. This happens because the search of each type of agglomeration occurs independently. Thus, a design problem might be related to more than one agglomeration.

Syntactic Agglomerations Suffice to Indicate Several Design Problems. Table 4 indicates that inter-component agglomerations are helpful for locating more instances of design problems than intra-component agglomerations. In four (out of seven) systems, at least almost half (45%) of the inter-component agglomerations were related to design problems. However, Table 4 also reveals that both types of syntactic agglomerations are useful to locate different de-

Table 4: Relation of Agglomerations and Design Problems: Precision and Recall

Agglom.	Intra-component			Inter-component			Syntactic			Semantic		
System	P	R	DP	P	R	DP	P	R	DP (*)	P	R	DP (**)
OODT	62%	30%	196	73%	83%	549	70%	97%	640 (535)	91%	65%	431 (201)
MM	42%	31%	12	82%	23%	9	57%	55%	21 (21)	100%	13%	5 (3)
HW	39%	11%	22	45%	87%	163	44%	87%	163 (141)	75%	100%	187 (24)
S1	43%	23%	64	51%	21%	58	47%	37%	104 (86)	81%	12%	34 (18)
S2	38%	9%	77	40%	7%	60	39%	13%	110 (83)	82%	5%	47 (21)
S3	38%	9%	76	39%	6%	58	38%	12%	109 (84)	75%	6%	50 (25)
S4	66%	89%	66	34%	60%	45	49%	93%	69 (27)	100%	13%	10 (6)
Median	42%	23%		45%	23%		47%	55%		82%	13%	
SD	0.118	0.281		0.182	0.349		0.109	0.368		0.107	0.370	

* Number of design problems exclusively related either to intra-component or to inter-component agglomerations.

** Number of design problems exclusively related to semantic agglomerations.

sign problems in most systems. As we expected, these two types of syntactic agglomerations are complementary. This finding can be observed by comparing the numbers of the sub-columns DP from the second to the fourth columns. In the column Syntactic, there are additional numbers between parentheses close to the numbers of design problems (sub-column DP). These numbers represent the amount of design problems that were exclusively related to either intra-component or inter-component agglomerations. Therefore, the subsets of design problems related to intra- and inter-component agglomerations in six systems are significantly different. The only exception was the Health Watcher. The reason is that all the design problems found in it were related to the communication between two components.

Semantic Agglomerations are Consistent Indicators. Semantic agglomerations were the most consistent indicators of design problems. This finding is based on the fact that semantic agglomerations were a good indicator of design problems across all the systems. The 5th column (ORs) of Table 3 shows that the chance of each anomaly within a semantic agglomeration being related to a design problem is more than seven times higher than each “detached” anomaly. This likelihood is higher than anomalies within syntactic agglomerations. In addition, semantic agglomerations presented the highest correlation with design problems, when compared to all other syntactic agglomerations. For example, 91% of the semantic agglomerations in OODT were related to 431 instances of design problems.

Semantic agglomerations are also often related to design problems even for the other systems where the absolute number of related design problems is lower: (i) the percentage of true positives was high (from 75% to 100%), and (ii) the percentage of false positives was low (from 0% to 25%). Considering the data from all systems, we observed a median of 82% of semantic agglomerations related to design problems. The standard deviation of 10.74% is low as well. That is, the percentage of semantic agglomerations related to design problems was high in all the target systems. Regarding the absolute number of design problems (DP), the semantic agglomerations were related to a high number of design problems in all the systems.

Syntactic and Semantic Agglomerations are Complementary. Considering each type of agglomeration, recall values were below 60% for most of the target systems (Table 4). Our results suggests that this occurs because syntactic and semantic agglomerations are complementary. In the last column of Table 4 (DP), there are additional

numbers between parentheses close to the numbers of design problems. These numbers represent the amount of design problems exclusively related to semantic agglomerations, i.e., design problems that could not be located only with syntactic agglomerations. Figure 2 shows an example of design problem located only in the context of semantic agglomeration.

Approximately 50% of the design problems related to semantic agglomerations have no relation to syntactic agglomerations in most of the systems. These design problems were often cases of *Scattered Concerns*, *Component Overload*, *Overused Interface* and *Delegating Abstraction*. This result suggests that semantic agglomerations are useful to complete the location of several design problems. These design problems are also hard to be located by developers as there is no syntactic relationship among the anomalous elements forming the agglomerations. However, existing work only focuses on characterizing code anomalies based on their syntactic relationships [1, 2, 30, 36]. Moreover, these design problems cannot be detected through the use of only concern location techniques. Several crosscutting concerns in the analyzed systems are not harmful, i.e., their implementations did not contain code anomalies and they were not actual sources of design problems.

Design Problems Indirectly Related to Agglomerations. We observed occurrences of design problems affecting methods or classes that are indirectly related to agglomerations even not participating of one. We observed three recurrent patterns of indirect relation: two for methods and one for classes. The first occurs when a method is implemented inside a class that is part of an agglomeration. The second is observed when a method is implemented in a class that have another method in an agglomeration. Finally, the third occurs when a class have one or more methods in agglomerations. We expect an increase in the recall values because, considering indirect relation, makes the correlation rule less strict. For instance, in S3 the recall value for semantic agglomerations increases from 6% to 51%. This suggests that an agglomeration may be helpful to identify design problems even without a direct relation. This becomes possible when developers analyze the full context of agglomerations, instead of analyzing anomalies individually.

Reducing False Positives of Syntactic Agglomerations. Even though syntactic agglomerations are often related to design problems, developers would still have to inspect and discard several irrelevant agglomerations. Table 4 shows that the number of false positives is higher than

60% in two systems. However, we observed that the combined use of semantic and syntactic agglomerations would significantly reduce the number of false positives yielded by the use of only syntactic agglomerations. If we only consider syntactic agglomerations that intersect with semantic agglomerations, the number of false positives for syntactic agglomerations would be reduced to 21% or less in all the systems. There is an intersection between a syntactic agglomeration X and a semantic agglomeration Y when at least one anomalous element take part of both X and Y. In summary, the joint use of syntactic and semantic agglomerations significantly enhances the location of design problems.

Full Automation of Semantic Agglomerations? The computation of semantic agglomerations relies on mapping of system concerns to indicate design problems. Tables 3 and 4 show the results for semantic agglomerations computed with concern mappings produced by the systems’ developers. However, it is important to check whether semantic agglomerations can also improve the location of design problems when concern mappings are automatically generated. As mentioned in Section 3.5, we also relied on concern mappings provided by Mallet in order to generate a second set of semantic agglomerations. Comparing to the developers’ mapping, we noticed Mallet generates longer lists of concerns and concern mappings. As a consequence, the use of Mallet results in a higher number of semantic agglomerations, which was related to more design problems than our original computation (Table 4). Semantic agglomerations generated with Mallet’s output lead to an average increase of 39.58% in the identification of design problems when compared to semantic agglomerations generated with developers’ mappings. The Mallet configuration and all the detailed results are available in our supplementary material [31].

4.3 Agglomerations as Congenital Problems?

This section addresses the RQ2: “What proportion of design problems early manifest themselves as agglomerations?” Early manifestation means the design problems were present in the first versions of a system. Therefore, they are likely to represent “congenital” design problems, i.e., they were introduced at design time. In order to answer this question, we computed for all the systems: (i) the number of design problems in the first versions of our target systems, and (ii) the proportion of design problems related to agglomerations.

The analysis of all initial versions revealed that a considerable number of design problems was introduced in the first version of each target system. As opposed to our expectation, all the initial versions presented a high correlation between agglomerations and design problems. We expected most of the design problems would be “evolutionary”, i.e., they would be introduced during the system evolution. However, the proportion of design problems related to agglomerations in early versions was approximately 75% for MM, S1, S2 and S4. The proportion for the other systems was close to 65%. We also observed that a considerable proportion of design problems are exclusively related to either syntactic or semantic agglomerations already in the first version.

These results show that change history analysis [11, 35, 41, 42] would not be an effective solution to reveal many instances of design problems. Many design problems are congenital and they are already “born” as agglomerations. It would be harder to remove them in later versions, when eventually the agglomeration’s anomalous elements start to

suffer co-changes through the later versions. An example of this case is the *Fat Interface* in the *Versioning* component in OODT (Section 2.2). It would be cumbersome to restructure the *Versioner* interface in later releases. The number of client classes for this interface increased from one (in the first release) to more than fifteen (last release).

Comparing Early and Late Versions. We also compared the nature of design problems and agglomerations in “early” and “late” versions of MM and HW systems. The comparison of the different versions of these two systems serves to illustrate most of our findings. For instance, let us consider the versions 1 and 8 of the MM system. Version 1 of MM (0.8 KSLOC) is smaller than version 8 (10 KSLOC). Therefore, the number of code anomalies and agglomerations is proportionally smaller in version 1. Even with fewer code elements, the first version already contained 23% of the agglomerations present in version 8. One of them, for example, is related to the *BaseController* class. In the first version, this class was identified as part of an intra-component agglomeration. That is, this anomalous class was related to another anomalous code element of the same component. This agglomeration was related to the *Component Overload* problem. In the subsequent versions, this agglomeration “expanded” to several code elements, including those located in other components. More specifically, in version 8, we identified that the same agglomeration became an inter-component agglomeration, involving classes that use the *BaseController* class. The inter-component agglomeration was related to emerging instances of *Fat Interface* and *Overused Interface* problems.

For the HW system, we compared versions 1 and 10. In this case, version 1 (6 KSLOC) was only somewhat smaller than version 10 (8 KSLOC). However, as in the case of the MM system, we observed again that some agglomerations found in a late version had already started to form in the first version. The first version already contained 39% of syntactic and semantic agglomerations found in version 10. All these observations suggest that agglomerations are effective to assist the identification and removal of design problems in the early versions of a system, which can prevent the introduction of more severe problems in subsequent versions.

5. RELATED WORK

Do Individual Anomalies Suffice? Tufano et al. [39] investigated when and why code smells are introduced in the context of evolving systems. Kim et al. [18], Lozano et al. [21] and Olbrich et al. [32] investigated the effect of individual code anomalies throughout the system’s evolution. These studies analyzed whether the number of code anomalies tend to increase over time, and how often they result in code changes. There is also a large body of work aiming at assessing the impact of individual occurrences of code anomalies on certain software quality attributes. For instance, Khomh et al. [16, 17] studied the relation between code anomalies and software change proneness. They concluded that anomalous classes tend to be more change-prone. In addition, corroborating with the work of Li and Shatnawi [20] and D’Ambros et al. [5], they also observed that classes infected by code smells tend to be more fault-prone than other classes [17]. Jürgens et al. [13] and Kapser et al. [14] found that one specific code anomaly adversely impacts on software maintenance. Sjöberg et al. [36] showed that single code anomalies were not related to maintenance effort.

Therefore, some of the results of these studies are controversial regarding the usefulness of individual code anomalies as indicators of software maintenance effort. However, none of the aforementioned studies have investigated the relation between anomaly agglomerations and design problems.

Agglomerations and Design Problems. While many authors investigated the impact of individual code anomalies, a few studies have focused on inter-related code anomalies. Abbes et al. [1] brought up the notion of interacting code anomalies and they studied its effects. They concluded that classes and methods identified as *God Classes* and *God Methods* in isolation had no effect on maintenance effort, but when appearing together, they led to a statistically significant increase in maintenance effort. Yamashita and Moonen [44] observed that inter-smell relationships negatively affect the maintenance. Macia [2] cataloged a set of patterns of inter-related code anomalies. However, none of these authors studied the impact of inter-related code anomalies on design problems, and none of them characterized or explored semantic relationships between code anomalies.

Architecturally connected problems. There are some studies that explored architecturally connected files to spot design problems in software systems. Kazman et al. [15] explored architecturally connected files to identify and quantify architectural debts in large-scale systems. Similarly to us, the authors identified design problems in the systems under analysis. However, we focus on the relation among code anomalies rather than the architectural connection among files. Mo et al. [29] proposed the detection of recurring architectural problems (called hotspot patterns) by the combination of history and architecture information. The hotspot patterns identify error- and change-prone files, and pinpoint specific architecture problems that may be the root causes of bugs and changes. The author's strategy to identify design problems is also based on structural dependencies among files. However, we showed that some design problems cannot be revealed based on structural dependencies only. They require to taken into account the semantic relation among anomalous code elements during their identification (Section 3.3). In addition, the approaches mentioned above rely on system's revision history. Consequently, they cannot be applied to early versions of the systems under analysis.

6. THREATS TO VALIDITY

Construct Validity. The major risk here is related to errors in the identification of design problems and code anomalies. The original developers reported the lists of design problems they found to be relevant in their systems. Therefore, these lists could be inaccurate and include irrelevant design problems. To mitigate this threat, we recruited developers with extensive knowledge of their system's design and with previous experience on design reviews and source code inspections. Moreover, we solicited a brief explanation about the severity of the reported design problems. Whenever we could not understand the nature and relevance of a design problem, we asked for further clarification. We discarded design problems that were not properly explained. Regarding the code anomalies, we selected detection strategies and thresholds that presented satisfactory results in a previous study (Section 3.5). In any case, if these strategies resulted in false positives or false negatives, they would have a similar effect on the computation of all the agglomerations.

Conclusion Validity. The main threat for conclusion

validity is the number of evaluated versions of each system. A study involving several versions of each system is always desired. However, it was impracticable in our study due to the number of systems (7) and the limited availability that original developers and architects had to participate in the study. We tried to mitigate this threat by selecting, for each different system, versions in a different lifecycle stage.

Internal and External Validity. The threats here concern the degree to which the findings can be generalized to the wider classes of subjects. In the experiment reported upon here, these threats are somewhat mitigated by the fact that we used systems with different sizes, purposes (academic, commercial and open-source) and domains, that were implemented using different design styles and that suffered from a different set of design problems. Furthermore, the target systems were developed by teams of different sizes and with different levels of software development skills.

7. CONCLUDING REMARKS

Recent empirical studies [23, 24, 25] suggest that individual code anomalies do not suffice to identify design problems. This finding represents a challenge for developers because they often need to spot design problems exclusively based on source code analysis, especially because of the lack of up-to-date design documents [24]. In order to decide whether and where a relevant design problem is prevailing in the program, programmers first need to know: (i) which are the code anomalies realizing the design problem, (ii) how these code anomalies are related to each other, and (iii) how these anomaly relationships are connected to the design problem. The gathering of all this scattered knowledge is time-consuming and error-prone.

In this work, we analyzed to what extent code anomaly agglomerations can help developers to locate design problems. We confirmed that code anomalies often “flock together” in order to embody a design problem. We studied the relation of design problems with both syntactic and semantic agglomerations in a longitudinal multi-case study involving 7 systems. Syntactic agglomerations are relevant indicators of design problems and help to discard an average of 600 irrelevant anomalies per system. We also found the combined use of syntactic and semantic agglomerations represents a more effective approach for locating design problems. Many congenital and evolutionary design problems were only related to semantic agglomerations. Moreover, a considerable proportion of design problems are congenital and manifest as agglomerations in the implementation. This result means that state-of-the-art solutions based on change history analysis are not adequate to promptly detect such problems. We plan to perform further investigations with the twofold goal of: (i) identifying efficient criteria for ranking code anomaly agglomerations in order to prioritize critical design problems, and (ii) conceiving a technique to analysis how anomaly agglomerations relate to each other throughout the software system histories.

8. ACKNOWLEDGMENTS

We thank professor Carlos José P. de Lucena and professor Nenad Medvidović for the support during this research. The research reported here was supported by the Capes Foundation, Ministry of Education of Brazil, FAPERJ and CNPq.

9. REFERENCES

- [1] M. Abbes, F. Khomh, Y.-G. Guéhéneuc, and G. Antoniol. An empirical study of the impact of two antipatterns, blob and spaghetti code, on program comprehension. In *Software Maintenance and Reengineering (CSMR), 2011 15th European Conference on*, pages 181–190, March 2011.
- [2] I. M. Bertrán. *On the Detection of Architecturally-Relevant Code Anomalies in Software Systems*. PhD thesis, Pontifical Catholic University of Rio de Janeiro (PUC-Rio), Rio de Janeiro, Brazil, 2013.
- [3] V. Bloomfield. *Using R for Numerical Analysis in Science and Engineering*. Chapman & Hall/CRC The R Series. Taylor & Francis, 2014.
- [4] J. Cornfield. A method of estimating comparative rates from clinical data; applications to cancer of the lung, breast, and cervix. *Journal of the National Cancer Institute*, 11(6):1269–1275, June 1951.
- [5] M. D’Ambros, A. Bacchelli, and M. Lanza. On the impact of design flaws on software defects. In *Quality Software (QSIC), 2010 10th International Conference on*, pages 23–31, July 2010.
- [6] R. A. Fisher. On the interpretation of χ^2 from contingency tables, and the calculation of p. *Journal of the Royal Statistical Society*, 85(1):pp. 87–94, 1922.
- [7] M. Fowler. *Refactoring: Improving the Design of Existing Code*. Addison-Wesley, Boston, MA, USA, 1999.
- [8] S. G. Ganesh, T. Sharma, and G. Suryanarayana. Towards a principle-based classification of structural design smells. *Journal of Object Technology*, pages 1–29, 2013.
- [9] J. Garcia, I. Ivkovic, and N. Medvidovic. A comparative analysis of software architecture recovery techniques. In *Automated Software Engineering (ASE), 2013 IEEE/ACM 28th International Conference on*, pages 486–496, Nov 2013.
- [10] J. Garcia, D. Popescu, G. Edwards, and N. Medvidovic. Identifying architectural bad smells. In *Software Maintenance and Reengineering, 2009. CSMR '09. 13th European Conference on*, pages 255–258, March 2009.
- [11] T. Gırba, S. Ducasse, A. Kuhn, R. Marinescu, and R. Daniel. Using concept analysis to detect co-change patterns. In *Ninth International Workshop on Principles of Software Evolution: In Conjunction with the 6th ESEC/FSE Joint Meeting, IWPSE '07*, pages 83–89, New York, NY, USA, 2007. ACM.
- [12] M. W. Godfrey and E. H. S. Lee. Secrets from the monster: Extracting Mozilla’s software architecture. In *Proc. of the Second Intl. Symposium on Constructing Software Engineering Tools (CoSET-00)*, June 2000.
- [13] E. Juergens, F. Deissenboeck, B. Hummel, and S. Wagner. Do code clones matter? In *Software Engineering, 2009. ICSE 2009. IEEE 31st International Conference on*, pages 485–495, May 2009.
- [14] C. J. Kapser and M. W. Godfrey. Cloning considered harmful? considered harmful: Patterns of cloning in software. *Empirical Softw. Engg.*, 13(6):645–692, Dec. 2008.
- [15] R. Kazman, Y. Cai, R. Mo, Q. Feng, L. Xiao, S. Haziyevev, V. Fedak, and A. Shapochka. A case study in locating the architectural roots of technical debt. In *Proceedings of the 37th International Conference on Software Engineering - Volume 2, ICSE '15*, pages 179–188, Piscataway, NJ, USA, 2015. IEEE Press.
- [16] F. Khomh, M. Di Penta, and Y. Guéhéneuc. An exploratory study of the impact of code smells on software change-proneness. In *Reverse Engineering, 2009. WCRE '09. 16th Working Conference on*, pages 75–84, Oct 2009.
- [17] F. Khomh, M. D. Penta, Y.-G. Guéhéneuc, and G. Antoniol. An exploratory study of the impact of antipatterns on class change- and fault-proneness. *Empirical Software Engineering*, 17(3):243–275, 2012.
- [18] M. Kim, V. Sazawal, D. Notkin, and G. Murphy. An empirical study of code clone genealogies. In *Proceedings of the 10th European Software Engineering Conference Held Jointly with 13th ACM SIGSOFT International Symposium on Foundations of Software Engineering, ESEC/FSE-13*, pages 187–196, New York, NY, USA, 2005. ACM.
- [19] M. Lanza, R. Marinescu, and S. Ducasse. *Object-Oriented Metrics in Practice*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2005.
- [20] W. Li and R. Shatnawi. An empirical study of the bad smells and class error probability in the post-release object-oriented system evolution. *Journal of Systems and Software*, 80(7):1120–1128, 2007.
- [21] A. Lozano and M. Wermelinger. Assessing the effect of clones on changeability. In *Software Maintenance, 2008. ICSM 2008. IEEE International Conference on*, pages 227–236, Sept 2008.
- [22] A. MacCormack, J. Rusnak, and C. Y. Baldwin. Exploring the structure of complex software designs: An empirical study of open source and proprietary code. *Manage. Sci.*, 52(7):1015–1030, July 2006.
- [23] I. Macia, R. Arcoverde, E. Cirilo, A. Garcia, and A. von Staa. Supporting the identification of architecturally-relevant code anomalies. In *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*, pages 662–665, Sept 2012.
- [24] I. Macia, R. Arcoverde, A. Garcia, C. Chavez, and A. von Staa. On the relevance of code anomalies for identifying architecture degradation symptoms. In *Software Maintenance and Reengineering (CSMR), 2012 16th European Conference on*, pages 277–286, March 2012.
- [25] I. Macia, J. Garcia, D. Popescu, A. Garcia, N. Medvidovic, and A. von Staa. Are automatically-detected code anomalies relevant to architectural modularity?: An exploratory analysis of evolving systems. In *Proceedings of the 11th Annual International Conference on Aspect-oriented Software Development, AOSD '12*, pages 167–178, New York, NY, USA, 2012. ACM.
- [26] R. C. Martin. *Agile Software Development: Principles, Patterns, and Practices*. Prentice Hall PTR, Upper Saddle River, NJ, USA, 2003.
- [27] C. A. Mattmann, D. J. Crichton, N. Medvidovic, and S. Hughes. A software architecture-based framework for highly distributed and data intensive scientific

- applications. In *Proceedings of the 28th International Conference on Software Engineering*, ICSE '06, pages 721–730, New York, NY, USA, 2006. ACM.
- [28] A. K. McCallum. Mallet: A machine learning for language toolkit. <http://mallet.cs.umass.edu>, 2002.
- [29] R. Mo, Y. Cai, R. Kazman, and L. Xiao. Hotspot patterns: The formal definition and automatic detection of architecture smells. In *Software Architecture (WICSA), 2015 12th Working IEEE/IFIP Conference on*, pages 51–60, May 2015.
- [30] N. Moha, Y. Guéhéneuc, L. Duchien, and A. Le Meur. Decor: A method for the specification and detection of code and design smells. *Software Engineering, IEEE Transactions on*, 36(1):20–36, Jan 2010.
- [31] W. Oizumi. Opus research group: Supplementary material, 2015. Availabel at <http://www.les.inf.puc-rio.br/opus/anomaliesFlockTogether>.
- [32] S. Olbrich, D. Cruzes, and D. I. Sjöberg. Are all code smells harmful? a study of god classes and brain classes in the evolution of three open source systems. In *Software Maintenance (ICSM), 2010 IEEE International Conference on*, pages 1–10, Sept 2010.
- [33] D. E. Perry and A. L. Wolf. Foundations for the study of software architecture. *SIGSOFT Softw. Eng. Notes*, 17(4):40–52, Oct. 1992.
- [34] S. Schach, B. Jin, D. Wright, G. Heller, and A. Offutt. Maintainability of the linux kernel. *Software, IEEE Proceedings* -, 149(1):18–23, Feb 2002.
- [35] R. Schwanke, L. Xiao, and Y. Cai. Measuring architecture quality by structure plus history analysis. In *Proceedings of the 2013 International Conference on Software Engineering*, ICSE '13, pages 891–900, Piscataway, NJ, USA, 2013. IEEE Press.
- [36] D. Sjöberg, A. Yamashita, B. Anda, A. Mockus, and T. Dyba. Quantifying the effect of code smells on maintenance effort. *Software Engineering, IEEE Transactions on*, 39(8):1144–1156, Aug 2013.
- [37] S. Soares, E. Laureano, and P. Borba. Implementing distribution and persistence aspects with aspectj. In *Proceedings of the 17th ACM SIGPLAN Conference on Object-oriented Programming, Systems, Languages, and Applications*, OOPSLA '02, pages 174–190, New York, NY, USA, 2002. ACM.
- [38] G. Suryanarayana, G. Samarthayam, and T. Sharma. *Refactoring for Software Design Smells: Managing Technical Debt*. Morgan Kaufmann, 2014.
- [39] M. Tufano, F. Palomba, G. Bavota, R. Oliveto, M. Di Penta, A. De Lucia, and D. Poshyvanyk. When and why your code starts to smell bad. In *Proceedings of the 37th International Conference on Software Engineering*, ICSE '15, New York, NY, USA, 2015. ACM.
- [40] J. van Gurp and J. Bosch. Design erosion: problems and causes. *Journal of Systems and Software*, 61(2):105 – 119, 2002.
- [41] S. Wong, Y. Cai, M. Kim, and M. Dalton. Detecting software modularity violations. In *Proceedings of the 33rd International Conference on Software Engineering*, ICSE '11, pages 411–420, New York, NY, USA, 2011. ACM.
- [42] L. Xiao, Y. Cai, and R. Kazman. Titan: A toolset that connects software architecture with quality analysis. In *Proceedings of the 22Nd ACM SIGSOFT International Symposium on Foundations of Software Engineering*, FSE 2014, pages 763–766, New York, NY, USA, 2014. ACM.
- [43] A. Yamashita and L. Moonen. Do code smells reflect important maintainability aspects? In *Software Maintenance (ICSM), 2012 28th IEEE International Conference on*, pages 306–315, Sept 2012.
- [44] A. Yamashita and L. Moonen. Exploring the impact of inter-smell relations on software maintainability: An empirical study. In *Software Engineering (ICSE), 2013 35th International Conference on*, pages 682–691, May 2013.
- [45] T. J. Young. Using aspectj to build a software product line for mobile devices. Master’s thesis, University of British Columbia, 2015.